

```
In [1]: # import to load dataset
import os
from torchvision import datasets, transforms
from torch.utils.data import random_split, DataLoader
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
import torch.optim as optim
from numba import cuda
use_cuda = torch.cuda.is_available()
import torchvision.transforms as T
import torchvision.datasets as D
import random
from torch.utils.data import Subset
```

```
In [2]: import os

large_dataset = ['CNV', 'DME', 'DRUSEN', 'NORMAL']
large_dataset_dir = '/lambda/nfs/resnet-filesystem/OCT_by_class'
# only count these as "images":
valid_exts = ('.png', '.jpg', '.jpeg')

def num_images(dataset_dir, folders):
    print(f"Number of images in each folder:")
    for folder in folders:
        path = os.path.join(dataset_dir, folder)
        if os.path.isdir(path):
            # only count files whose name ends with one of valid_exts
            num_files = sum(
                1 for fname in os.listdir(path)
                if fname.lower().endswith(valid_exts)
            )
            print(f"{folder}: {num_files}")
        else:
            print(f"Folder '{folder}' does not exist in the dataset directory.")

num_images(large_dataset_dir, large_dataset)
```

Number of images in each folder:

CNV: 3000
DME: 3000
DRUSEN: 3000
NORMAL: 3000

```
In [3]: from torchvision import transforms as T
from torchvision.datasets import ImageFolder
from torch.utils.data import random_split, DataLoader
```

```
In [4]: # import to load dataset
from torchvision.datasets import ImageFolder
from torch.utils.data import random_split, DataLoader
import torch

transform = T.ToTensor()
dataset = ImageFolder(root=large_dataset_dir, transform=transform)

# split the data: 70% training, 15% validation, 15% testing
total_len = len(dataset)
train_len = int(0.7 * total_len)
val_len = int(0.15 * total_len)
test_len = total_len - train_len - val_len
train_data, val_data, test_data = random_split(dataset, [train_len, val_len, test_len], generator=torch.Generator().manual_seed(42))

# define dataloader parameters
batch_size = 32

# prepare data loaders
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_data, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False)

# check the number of training, validation, and test images alongside the percentage of training, validation, and testing (check)
print(f"Number of training images: {len(train_data)} Percent: {100 * len(train_data)/total_len:.2f}")
print(f"Number of validation images: {len(val_data)} Percent: {100 * len(val_data)/total_len:.2f}")
print(f"Number of test images: {len(test_data)} Percent: {100 * len(test_data)/total_len:.2f}")
```

Number of training images: 8400 Percent: 70.00
Number of validation images: 1800 Percent: 15.00
Number of test images: 1800 Percent: 15.00

```
In [5]: import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import models, transforms
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
```

```
In [6]: def get_accuracy(model, data_loader):
    correct = 0
    total = 0
    for imgs, labels in data_loader:

        #####
        #To Enable GPU Usage
        if use_cuda and torch.cuda.is_available():
            imgs = imgs.cuda()
            labels = labels.cuda()
        #####
```

```

output = model(imgs)

# select index with maximum prediction score
pred = output.max(1, keepdim=True)[1]
correct += pred.eq(labels.view_as(pred)).sum().item()
total += imgs.shape[0]
return correct / total

```

```

In [7]: from tqdm import tqdm
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt

def train(model, train_data, val_data, batch_size=64, learning_rate = 0.001, num_epochs=20):
    from torch.utils.data import DataLoader

    train_loader = DataLoader(
        train_data,
        batch_size=batch_size,
        shuffle=True,
        num_workers=4,
        pin_memory=True,
        prefetch_factor=2
    )

    val_loader = DataLoader(
        val_data,
        batch_size=batch_size,
        shuffle=False,
        num_workers=4,
        pin_memory=True
    )

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)

    iters, losses, train_acc, val_acc = [], [], [], []

    # training
    n = 0 # the number of iterations
    for epoch in range(num_epochs):
        for imgs, labels in tqdm(train_loader, desc="Training", leave=False):

            #####
            #To Enable GPU Usage
            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()
            #####

            out = model(imgs) # forward pass
            loss = criterion(out, labels) # compute the total loss
            loss.backward() # backward pass (compute parameter updates)
            optimizer.step() # make the updates for each parameter
            optimizer.zero_grad() # a clean up step for PyTorch

            # save the current training information
            iters.append(n)
            losses.append(float(loss)/batch_size) # compute *average* Loss
            train_acc.append(get_accuracy(model, train_loader)) # compute training accuracy
            val_acc.append(get_accuracy(model, val_loader)) # compute validation accuracy
            n += 1

        print(f"Epoch {epoch+1}: Train acc: {train_acc[-1]:.4f} | Validation acc: {val_acc[-1]:.4f}")
        torch.cuda.empty_cache()
        #model_path = get_model_name(model.name, batch_size, learning_rate, epoch)
        #torch.save(model.state_dict(), model_path)

    # plotting
    plt.title("Training Curve")
    plt.plot(iters, losses, label="Train")
    plt.xlabel("Iterations")
    plt.ylabel("Loss")
    plt.show()

    plt.title("Training Curve")
    plt.plot(iters, train_acc, label="Train")
    plt.plot(iters, val_acc, label="Validation")
    plt.xlabel("Iterations")
    plt.ylabel("Training Accuracy")
    plt.legend(loc='best')
    plt.show()

    print("Final Training Accuracy: {}".format(train_acc[-1]))
    print("Final Validation Accuracy: {}".format(val_acc[-1]))

```

```

In [8]: from torchvision.models import ResNet50_Weights
resnet50 = models.resnet50(weights=ResNet50_Weights.IMAGENET1K_V1)

# freeze weights
for param in resnet50.parameters():
    param.requires_grad = False

in_features = resnet50.fc.in_features
resnet50.fc = nn.Linear(in_features, 4)

```

```

In [10]: use_cuda = True

if use_cuda and torch.cuda.is_available():
    resnet50 = resnet50.to('cuda:0')

```

```

device = 'cuda:0'
print('CUDA is available! Training on GPU ...')
else:
    device = 'cpu'
    print('CUDA is not available. Training on CPU ...')

train(resnet50, train_data, val_data, batch_size=256, learning_rate = 0.002, num_epochs=5)

```

CUDA is available! Training on GPU ...



Epoch 1: Train acc: 0.6687 | Validation acc: 0.6700

Epoch 2: Train acc: 0.7154 | Validation acc: 0.7028

Epoch 3: Train acc: 0.7262 | Validation acc: 0.7150

Epoch 4: Train acc: 0.7374 | Validation acc: 0.7372

Epoch 5: Train acc: 0.7313 | Validation acc: 0.7211



Final Training Accuracy: 0.7313095238095239

Final Validation Accuracy: 0.7211111111111111

```

In [11]: # compute the test accuracy for resnet50
test_acc = get_accuracy(resnet50, test_loader)
print(f"Test accuracy: {test_acc:.4f}")

```

Test accuracy: 0.7200

```

In [12]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# confusion matrix
def plot_confusion_matrix(model, data_loader, class_names):
    # set model into evaluation mode
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for imgs, labels in data_loader:
            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()

            output = model(imgs)
            preds = output.argmax(dim=1)

            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    # compute the confusion matrix
    cm = confusion_matrix(all_labels, all_preds)

    # plot the confusion matrix
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
    disp.plot(cmap='Blues', values_format='d')
    plt.title("Confusion Matrix")
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title('Confusion Matrix')
    plt.show()


    return cm

```

```

In [13]: class_names = ['CNV', 'DME', 'Drusen', 'Normal']
plot_confusion_matrix(resnet50, test_loader, class_names)

```



```
Out[13]: array([[314, 23, 123, 6],
               [ 29, 295, 71, 31],
               [ 32, 17, 362, 47],
               [ 3, 31, 75, 341]])
```

```
In [14]: from sklearn.metrics import roc_curve, auc
         from sklearn.preprocessing import label_binarize

         # plot roc curve
         def plot_roc_curve(model, data_loader, class_names):
             model.eval()
             all_labels = []
             all_probs = []

             with torch.no_grad():
                 for imgs, labels in data_loader:
                     if use_cuda and torch.cuda.is_available():
                         imgs = imgs.cuda()
                         labels = labels.cuda()

                     output = model(imgs)
                     probs = torch.softmax(output, dim=1)

                     all_probs.extend(probs.cpu().numpy())
                     all_labels.extend(labels.cpu().numpy())

             all_probs = np.array(all_probs)
             all_labels = np.array(all_labels)

             # binarize the labels for multi-class ROC
             y_true = label_binarize(all_labels, classes=np.arange(len(class_names)))
             n_classes = y_true.shape[1]

             # compute ROC and AUC for each class
             fpr = dict()
             tpr = dict()
             roc_auc = dict()

             for i in range(n_classes):
                 fpr[i], tpr[i], _ = roc_curve(y_true[:, i], all_probs[:, i])
                 roc_auc[i] = auc(fpr[i], tpr[i])

             # plot the ROC curve
             plt.figure()
             for i in range(n_classes):
                 plt.plot(fpr[i], tpr[i], label=f"{class_names[i]} (AUC = {roc_auc[i]:.2f})")

             plt.plot([0, 1], [0, 1], 'k--') # random classifier
             plt.xlabel('False Positive Rate')
             plt.ylabel('True Positive Rate')
             plt.title('Receiver Operating Characteristic (ROC) Curve')
             plt.legend(loc='lower right')
             plt.grid(True)
             plt.show()

             return roc_auc
```

```
In [15]: roc_auc = plot_roc_curve(resnet50, test_loader, class_names)
```

