

```
In [1]: import os

large_dataset = ['CNV', 'DME', 'DRUSEN', 'NORMAL']
large_dataset_dir = '/lambda/nfs/resnet-filesystem/OCT_by_class' # change directory

def num_images(dir, folders):
    print(f"Number of images in each folder:")
    for folder in folders:
        path = os.path.join(dir, folder)
        if os.path.isdir(path):
            num_files = len(os.listdir(path))
            print(f"{folder}: {num_files}")
        else:
            print(f"Folder '{folder}' does not exist in the dataset directory.")

num_images(large_dataset_dir, large_dataset)
```

Number of images in each folder:

CNV: 3000
DME: 3000
DRUSEN: 3000
NORMAL: 3000

```
In [2]: from torchvision import transforms as T
from torchvision.datasets import ImageFolder
from torch.utils.data import random_split, DataLoader
```

```
In [3]: # import to load dataset
from torchvision.datasets import ImageFolder
from torch.utils.data import random_split, DataLoader
import torch

transform = T.ToTensor()
dataset = ImageFolder(root=large_dataset_dir, transform=transform)

# split the data: 70% training, 15% validation, 15% testing
total_len = len(dataset)
train_len = int(0.7 * total_len)
val_len = int(0.15 * total_len)
test_len = total_len - train_len - val_len
train_data, val_data, test_data = random_split(dataset, [train_len, val_len, test_len], generator=torch.Generator().manual_seed(42))

# define dataloader parameters
batch_size = 32

# prepare data loaders
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_data, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False)

# check the number of training, validation, and test images alongside the percentage of training, validation, and testing (check)
print(f"Number of training images: {len(train_data)} Percent: {100 * len(train_data)/total_len:.2f}")
print(f"Number of validation images: {len(val_data)} Percent: {100 * len(val_data)/total_len:.2f}")
print(f"Number of test images: {len(test_data)} Percent: {100 * len(test_data)/total_len:.2f}")
```

Number of training images: 8400 Percent: 70.00
Number of validation images: 1800 Percent: 15.00
Number of test images: 1800 Percent: 15.00

```
In [4]: def get_accuracy(model, data_loader):
    correct = 0
    total = 0
    for imgs, labels in data_loader:

        #####
        #To Enable GPU Usage
        if use_cuda and torch.cuda.is_available():
            imgs = imgs.cuda()
            labels = labels.cuda()
        #####

        output = model(imgs)

        # select index with maximum prediction score
        pred = output.max(1, keepdim=True)[1]
        correct += pred.eq(labels.view_as(pred)).sum().item()
        total += imgs.shape[0]
    return correct / total
```

```
In [5]: from tqdm import tqdm
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt

def train(model, train_data, val_data, batch_size=64, learning_rate = 0.001, num_epochs=20):
    from torch.utils.data import DataLoader

    # optimize the data loaders
    train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True, num_workers=4, pin_memory=True, prefetch_factor=2)
    val_loader = DataLoader(val_data, batch_size=batch_size, shuffle=False, num_workers=4, pin_memory=True)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)

    iters, losses, train_acc, val_acc = [], [], [], []

    # training
    n = 0 # the number of iterations
    for epoch in range(num_epochs):
        for imgs, labels in tqdm(train_loader, desc="Training", leave=False):
```

```
#####
#To Enable GPU Usage
if use_cuda and torch.cuda.is_available():
    imgs = imgs.cuda()
    labels = labels.cuda()
#####

out = model(imgs)          # forward pass
loss = criterion(out, labels) # compute the total Loss
loss.backward()            # backward pass (compute parameter updates)
optimizer.step()           # make the updates for each parameter
optimizer.zero_grad()      # a clean up step for PyTorch

# save the current training information
iters.append(n)
losses.append(float(loss)/batch_size) # compute *average* Loss
train_acc.append(get_accuracy(model, train_loader)) # compute training accuracy
val_acc.append(get_accuracy(model, val_loader)) # compute validation accuracy
n += 1

print(f"Epoch {epoch+1}: Train acc: {train_acc[-1]:.4f} | Validation acc: {val_acc[-1]:.4f}")
#model_path = get_model_name(model.name, batch_size, learning_rate, epoch)
#torch.save(model.state_dict(), model_path)

# plotting
plt.title("Training Curve")
plt.plot(iters, losses, label="Train")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()

plt.title("Training Curve")
plt.plot(iters, train_acc, label="Train")
plt.plot(iters, val_acc, label="Validation")
plt.xlabel("Iterations")
plt.ylabel("Training Accuracy")
plt.legend(loc='best')
plt.show()

print("Final Training Accuracy: {}".format(train_acc[-1]))
print("Final Validation Accuracy: {}".format(val_acc[-1]))
```

```
In [6]: import torchvision
vggModel = torchvision.models.vgg16(pretrained=True)
```

```
/usr/lib/python3/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/lib/python3/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=VGG16_Weights.IMAGENET1K_V1`. You can also use `weights=VGG16_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```
In [7]: vggModel.classifier[6] = nn.Linear(in_features=4096, out_features=4)
```

```
In [8]: # freeze all feature extractor layers
for param in vggModel.features.parameters():
    param.requires_grad = False
```

```
In [9]: use_cuda = True

if use_cuda and torch.cuda.is_available():
    vggModel = vggModel.to('cuda:0')
    device = 'cuda:0'
    print('CUDA is available! Training on GPU ...')
else:
    device = 'cpu'
    print('CUDA is not available. Training on CPU ...')

train(vggModel, train_data, val_data, batch_size=256, learning_rate = 0.002, num_epochs=5)
```

CUDA is available! Training on GPU ...

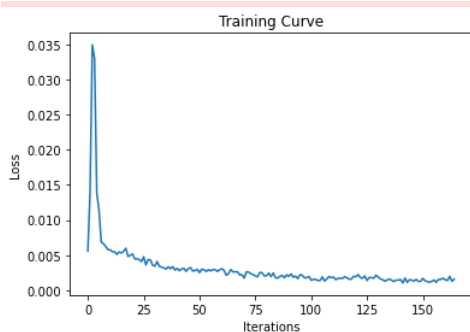
Epoch 1: Train acc: 0.6179 | Validation acc: 0.6400

Epoch 2: Train acc: 0.7831 | Validation acc: 0.7489

Epoch 3: Train acc: 0.8304 | Validation acc: 0.7728

Epoch 4: Train acc: 0.8560 | Validation acc: 0.7772

Epoch 5: Train acc: 0.8664 | Validation acc: 0.7761





Final Training Accuracy: 0.8664285714285714
Final Validation Accuracy: 0.7761111111111111

```
In [10]: # compute the test accuracy for vgg16
test_acc = get_accuracy(vggModel, test_loader)
print(f"Test accuracy: {test_acc:.4f}")

Test accuracy: 0.7811
```

```
In [11]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# confusion matrix
def plot_confusion_matrix(model, data_loader, class_names):
    # set model into evaluation mode
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for imgs, labels in data_loader:
            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()

            output = model(imgs)
            preds = output.argmax(dim=1)


            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    # compute the confusion matrix
    cm = confusion_matrix(all_labels, all_preds)

    # plot the confusion matrix
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
    disp.plot(cmap='Blues', values_format='d')
    plt.title("Confusion Matrix")
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title('Confusion Matrix')
    plt.show()

    return cm
```

```
In [12]: class_names = ['CNV', 'DME', 'Drusen', 'Normal']
plot_confusion_matrix(vggModel, test_loader, class_names)
```



```
Out[12]: array([[367, 10, 88, 1],
 [ 31, 329, 41, 25],
 [ 35, 8, 399, 16],
 [ 2, 17, 66, 365]])
```

```
In [13]: from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
import numpy as np

# plot roc curve
def plot_roc_curve(model, data_loader, class_names):
    model.eval()
    all_labels = []
    all_probs = []

    with torch.no_grad():
        for imgs, labels in data_loader:
            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()

            output = model(imgs)
            probs = torch.softmax(output, dim=1)
```

```

    all_probs.extend(probs.cpu().numpy())
    all_labels.extend(labels.cpu().numpy())

all_probs = np.array(all_probs)
all_labels = np.array(all_labels)

# binarize the labels for multi-class ROC
y_true = label_binarize(all_labels, classes=np.arange(len(class_names)))
n_classes = y_true.shape[1]

# compute ROC and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_true[:, i], all_probs[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# plot the ROC curve
plt.figure()
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f"{class_names[i]} (AUC = {roc_auc[i]:.2f})")

plt.plot([0, 1], [0, 1], 'k--') # random classifier
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

return roc_auc

```

In [14]: `roc_auc = plot_roc_curve(vggModel, test_loader, class_names)`

